

Name : Garimella Suresh Babu
Student Id: 700741367
User Id :SXG13670@ucmo.edu

GitHub Link:-

https://github.com/Suresh-Garimella/Principal_Component_Analysis-Linear_Discriminant_Analysis

Video Link:-

<https://github.com/Suresh-Garimella/Linear-Regression-Clustering/blob/main/ScreenRecording.mkv>

1. Principal Component Analysis

a. Apply PCA on CC dataset.

Importing all the essential libraries and packages to perform PCA on the given Dataset.

Namely Packages from Sklearn - to perform PCA.

Pandas - to Maintain Dataframes.

Matplotlib - to visualize the data.

```

from sklearn.model_selection import train_test_split
import pandas as pd
# from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')

```

```

In [86]: df= pd.read_csv(r"CC.csv")

df.head()

```

Out[86]:

| | CUST_ID | BALANCE | BALANCE_FREQUENCY | PURCHASES | ONEOFF_PURCHASES | INSTALLMENTS_PURCHASES |
|---|---------|-------------|-------------------|-----------|------------------|------------------------|
| 0 | C10001 | 40.900749 | 0.818182 | 95.40 | 0.00 | 95.4 |
| 1 | C10002 | 3202.467416 | 0.909091 | 0.00 | 0.00 | 0.0 |
| 2 | C10003 | 2495.148862 | 1.000000 | 773.17 | 773.17 | 0.0 |
| 3 | C10004 | 1666.670542 | 0.636364 | 1499.00 | 1499.00 | 0.0 |
| 4 | C10005 | 817.714335 | 1.000000 | 16.00 | 16.00 | 0.0 |

```

In [65]: df.shape

```

Out[65]: (8950, 18)

1. Apply Linear Regression to the provided dataset using underlying steps.

a. Import the given “Salary_Data.csv”

Imports : pandas , numpy and matplotlib libraries

Using the pandas library we read the Salary_Data.csv file.

After that we are reading the CC.csv file using pandas

```
In [85]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import pandas as pd
# from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

```
In [86]: df= pd.read_csv(r"CC.csv")
df.head()
```

```
Out[86]:
```

| | CUST_ID | BALANCE | BALANCE_FREQUENCY | PURCHASES | ONEOFF_PURCHASES | INSTALLMENTS_PURCHASES | CASH_ADVANCE | PURCHASES_FREQUEN |
|---|---------|-------------|-------------------|-----------|------------------|------------------------|--------------|-------------------|
| 0 | C10001 | 40.900749 | 0.818182 | 95.40 | 0.00 | 95.4 | 0.000000 | 0.166 |
| 1 | C10002 | 3202.467416 | 0.909091 | 0.00 | 0.00 | 0.0 | 6442.945483 | 0.000 |
| 2 | C10003 | 2495.148862 | 1.000000 | 773.17 | 773.17 | 0.0 | 0.000000 | 1.000 |
| 3 | C10004 | 1666.670542 | 0.636364 | 1499.00 | 1499.00 | 0.0 | 205.788017 | 0.083 |
| 4 | C10005 | 817.714335 | 1.000000 | 16.00 | 16.00 | 0.0 | 0.000000 | 0.083 |

```
In [65]: df.shape
```

```
Out[65]: (8950, 18)
```

Then we Found that the data contains null values. And we Replace the Null the null values with the mean of their column.

Then by preprocessing we deleted the CUST_ID column.

```
[68]: mean1=df['CREDIT_LIMIT'].mean()
      mean2=df['MINIMUM_PAYMENTS'].mean()
      df['CREDIT_LIMIT'].fillna(value=mean1, inplace=True)
      df['MINIMUM_PAYMENTS'].fillna(value=mean2, inplace=True)

In [ ]:

[69]: df.values

In [69]: array([[ 'C10001', 40.900749, 0.818182, ..., 139.50978700000002, 0.0, 12],
                [ 'C10002', 3202.467416, 0.909091, ..., 1072.340217,
                  0.22222199999999998, 12],
                [ 'C10003', 2495.148862, 1.0, ..., 627.284787, 0.0, 12],
                ...,
                [ 'C19188', 23.398673000000002, 0.833333, ..., 82.418369, 0.25, 6],
                [ 'C19189', 13.457564000000001, 0.833333, ..., 55.755628, 0.25, 6],
                [ 'C19190', 372.708075, 0.666667, ..., 88.288956, 0.0, 6]],
               dtype=object)

[70]: df['TENURE'].value_counts()

In [70]: 12    7584
          11    365
          10    236
           6    204
           8    196
           7    190
           9    175
          Name: TENURE, dtype: int64

[71]: del df['CUST_ID']

      df.head()
```

Then we need to Drop the Decision Attribute “TENURE” and store the rest of columns in X.And also store the decision attribute column in Y.

```

In [72]:
X = df.drop('TENURE',axis=1).values
print(X)
y = df['TENURE'].values
print(y)

[[4.09007490e+01 8.18182000e-01 9.54000000e+01 ... 2.01802084e+02
 1.39509787e+02 0.00000000e+00]
 [3.20246742e+03 9.09091000e-01 0.00000000e+00 ... 4.10303260e+03
 1.07234022e+03 2.22222000e-01]
 [2.49514886e+03 1.00000000e+00 7.73170000e+02 ... 6.22066742e+02
 6.27284787e+02 0.00000000e+00]
 ...
 [2.33986730e+01 8.33333000e-01 1.44400000e+02 ... 8.12707750e+01
 8.24183690e+01 2.50000000e-01]
 [1.34575640e+01 8.33333000e-01 0.00000000e+00 ... 5.25499590e+01
 5.57556280e+01 2.50000000e-01]
 [3.72708075e+02 6.66667000e-01 1.09325000e+03 ... 6.31654040e+01
 8.82889560e+01 0.00000000e+00]]
[12 12 12 ... 6 6 6]

```

Then transformed the data using PCA with 2 components that means the final dataset has only 2 columns excluding the final attribute. Then we added the Tenure attribute to the dataframe.

```

In [73]: %%time
pca2 = PCA(n_components=2)
principalComponents = pca2.fit_transform(X)

principalDf = pd.DataFrame(data = principalComponents, columns = ['principal component 1', 'principal component 2'])

finalDf = pd.concat([principalDf, df[['TENURE']]], axis = 1)
finalDf.head()

Wall time: 58.4 ms

```

```

Out[73]:


```

| | principal component 1 | principal component 2 | TENURE |
|---|-----------------------|-----------------------|--------|
| 0 | -4326.383979 | 921.566882 | 12 |
| 1 | 4118.916665 | -2432.846346 | 12 |
| 2 | 1497.907641 | -1997.578694 | 12 |
| 3 | 1394.548536 | -1488.743453 | 12 |
| 4 | -3743.351896 | 757.342657 | 12 |

Visualization of the final dataset after PCA.

```
In [74]: plt.figure(figsize=(7,7))
plt.scatter(finalDf['principal component 1'],finalDf['principal component 2'],c=df['TENURE'],cmap='prism', s =5)
plt.xlabel('pc1')
plt.ylabel('pc2')
```

Out[74]: Text(0, 0.5, 'pc2')




```
In [75]: %%time
```

```
##elbow method to know the number of clusters

from sklearn.cluster import KMeans
wcss = []
for i in range(1,11):
    kmeans = KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0)
    kmeans.fit(finalDf)
    wcss.append(kmeans.inertia_)

plt.plot(range(1,11),wcss)
plt.title('the elbow method')
plt.xlabel('Number of Clusters')
plt.ylabel('wcss')
plt.show()
```



Wall time: 9.27 s

Scaling + PCA + K-Means

```
In [77]: scaler = StandardScaler()  
X_Scale = scaler.fit_transform(X)
```

```
In [78]: %%time  
  
pca2 = PCA(n_components=2)  
principalComponents = pca2.fit_transform(X_Scale)  
  
principalDf = pd.DataFrame(data = principalComponents, columns = ['principal component 1', 'principal component 2'])  
  
finalDf = pd.concat([principalDf, df[['TENURE']]], axis = 1)  
finalDf.head()
```


Wall time: 44.9 ms

```
Out[78]:
```

| | principal component 1 | principal component 2 | TENURE |
|---|-----------------------|-----------------------|--------|
| 0 | -1.718890 | -1.072934 | 12 |
| 1 | -1.169302 | 2.509332 | 12 |
| 2 | 0.938413 | -0.382604 | 12 |
| 3 | -0.907501 | 0.045864 | 12 |
| 4 | -1.637828 | -0.684970 | 12 |

```
In [79]: plt.figure(figsize=(7,7))  
plt.scatter(finalDf['principal component 1'],finalDf['principal component 2'],c=df['TENURE'],cmap='prism', s =5)  
plt.xlabel('pc1')  
  
plt.ylabel('pc2')
```


Out[79]: Text(0, 0.5, 'pc2')



```
In [80]: %%time
        ##elbow method to know the number of clusters

        from sklearn.cluster import KMeans
        wcss = []
        for i in range(1,11):
            kmeans = KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10,random_state=0)
            kmeans.fit(finalDf)
            wcss.append(kmeans.inertia_)

        plt.plot(range(1,11),wcss)
        plt.title('the elbow method')
        plt.xlabel('Number of Clusters')
        plt.ylabel('Wcss')
        plt.show()
```



Wall time: 12.5 s


```
In [81]: # Calculate the silhouette score for the above clustering
#since the elbow point is at 3
nclusters = 3 # this is the k in kmeans
km = KMeans(n_clusters=nclusters)
km.fit(finalDf)

y_cluster_kmeans = km.predict(finalDf)
from sklearn import metrics
score = metrics.silhouette_score(finalDf, y_cluster_kmeans)
print(score)
```

0.38366768475212315

```
In [43]: df= pd.read_csv(r"pd_speech_features.csv")
df.head()
```

```
Out[43]:
```

| | id | gender | PPE | DFA | RPDE | numPulses | numPeriodsPulses | meanPeriodPulses | stdDevPeriodPulses | locPctJitter | ... | tqwt_kurtosisValue_dec_28 | tq |
|---|----|--------|---------|---------|---------|-----------|------------------|------------------|--------------------|--------------|-----|---------------------------|----|
| 0 | 0 | 1 | 0.85247 | 0.71826 | 0.57227 | 240 | 239 | 0.008064 | 0.000087 | 0.00218 | ... | 1.5620 | |
| 1 | 0 | 1 | 0.76686 | 0.69481 | 0.53966 | 234 | 233 | 0.008258 | 0.000073 | 0.00195 | ... | 1.5589 | |
| 2 | 0 | 1 | 0.85083 | 0.67604 | 0.58982 | 232 | 231 | 0.008340 | 0.000060 | 0.00176 | ... | 1.5643 | |
| 3 | 1 | 0 | 0.41121 | 0.79672 | 0.59257 | 178 | 177 | 0.010858 | 0.000183 | 0.00419 | ... | 3.7805 | |
| 4 | 1 | 0 | 0.32790 | 0.79782 | 0.53028 | 236 | 235 | 0.008162 | 0.002669 | 0.00535 | ... | 6.1727 | |

5 rows × 755 columns

```
n [44]: df.isnull().any()
```

```
ut[44]: id                False
gender                 False
PPE                   False
DFA                   False
RPDE                  False
...
tqwt_kurtosisValue_dec_33  False
tqwt_kurtosisValue_dec_34  False
tqwt_kurtosisValue_dec_35  False
tqwt_kurtosisValue_dec_36  False
class                 False
Length: 755, dtype: bool
```

```
n [45]: x = df.drop('class',axis=1).values
y = df['class'].values
```

```
n [46]: print(X)
```

```
[[ 0.      1.      0.85247 ... 2.6202    3.0004    18.9405 ]
 [ 0.      1.      0.76686 ... 6.5245    6.3431    45.178  ]
 [ 0.      1.      0.85083 ... 2.9199    3.1495     4.7666 ]
 ...
 [251.     0.      0.88389 ... 3.5377    3.3545     5.0424 ]
 [251.     0.      0.83782 ... 2.6801    2.8332     3.7131 ]
 [251.     0.      0.81304 ... 4.0116    2.6217     3.1527 ]]
```

```
In [47]: scaler = StandardScaler()
X_Scale = scaler.fit_transform(X)
X_Scale
```

```
Out[47]: array([[ -1.72519117,  0.96874225,  0.62764391, ..., -0.775137 ,
                  -0.81472704, -0.36659507],
                 [ -1.72519117,  0.96874225,  0.12161952, ..., -0.52664699,
                  -0.58297219,  0.40039616],
                 [ -1.72519117,  0.96874225,  0.61795018, ..., -0.75606253,
                  -0.8043897 , -0.7809355 ],
                 ...,
                 [ 1.72519117, -1.03226633,  0.81336154, ..., -0.71674252,
                  -0.79017671, -0.77287314],
                 [ 1.72519117, -1.03226633,  0.54105055, ..., -0.77132466,
                  -0.82631929, -0.81173208],
                 [ 1.72519117, -1.03226633,  0.3945807 , ..., -0.68658105,
                  -0.84098293, -0.82811405]])
```

```
In [48]: X_Scale.shape
```

```
Out[48]: (756, 754)
```

```

1 [49]: pca3 = PCA(n_components=3)
principalComponents = pca3.fit_transform(X_Scale)

principalDf = pd.DataFrame(data = principalComponents, columns = ['principal component 1', 'principal component 2', 'principal component 3'])

finalDf = pd.concat([principalDf, df[['class']]], axis = 1)
finalDf.head()

```

```

In[49]:

```

| | principal component 1 | principal component 2 | principal component 3 | class |
|---|-----------------------|-----------------------|-----------------------|-------|
| 0 | -10.047372 | 1.471076 | -6.846404 | 1 |
| 1 | -10.637725 | 1.583749 | -6.830977 | 1 |
| 2 | -13.516185 | -1.253542 | -6.818697 | 1 |
| 3 | -9.155083 | 8.833600 | 15.290898 | 1 |
| 4 | -6.764470 | 4.611465 | 15.637113 | 1 |

```

[50]: from mpl_toolkits.mplot3d import Axes3D

fig = plt.figure(figsize=(9,9))
axes = Axes3D(fig)
axes.set_title('PCA Representation', size=14)
axes.set_xlabel('PC1')
axes.set_ylabel('PC2')
axes.set_zlabel('PC3')


axes.scatter(finalDf['principal component 1'], finalDf['principal component 2'], finalDf['principal component 3'], c=finalDf['class'])

```

```

[50]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x24100d589c8>

```



Apply SVM

```
In [51]: from sklearn.model_selection import train_test_split, cross_validate
training_set = finalDf[[:-1]]
labels = finalDf["class"]
```

```
X_train, X_val, Y_train, Y_val = train_test_split(training_set, labels, test_size=0.3, random_state=0)
```

```
In [52]: X_train.head()
```

Out[52]:

| | principal component 1 | principal component 2 | principal component 3 | class |
|-----|-----------------------|-----------------------|-----------------------|-------|
| 175 | -1.888843 | -2.410033 | -6.888701 | 1 |
| 64 | -9.681345 | -2.826047 | -5.915025 | 1 |
| 440 | -4.096262 | 9.699011 | -0.957693 | 1 |
| 555 | -0.576323 | 10.768601 | 1.271319 | 1 |
| 436 | 9.454904 | 8.833950 | 0.601404 | 1 |

In [53]:

```
from sklearn.svm import SVC

classifier = LinearSVC()

classifier.fit(X_train, Y_train)

y_pred = classifier.predict(X_val)

from sklearn.metrics import classification_report
# Summary of the predictions made by the classifier
print(classification_report(Y_val, y_pred))

from sklearn.metrics import confusion_matrix
print(confusion_matrix(Y_val, y_pred))
# Package used to get the Accuracy score
from sklearn.metrics import accuracy_score
print('accuracy is', accuracy_score(Y_val, y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.00 | 0.00 | 0.00 | 57 |
| 1 | 0.75 | 0.99 | 0.85 | 170 |
| accuracy | | | 0.74 | 227 |
| macro avg | 0.37 | 0.50 | 0.43 | 227 |
| weighted avg | 0.56 | 0.74 | 0.64 | 227 |

```
[[ 0 57]
 [ 1 169]]
```

accuracy is 0.7444933920704846

```
n [76]: import pandas as pd
        from sklearn.metrics import accuracy_score
        import warnings
        warnings.filterwarnings('ignore')
```

```
n [77]: df= pd.read_csv(r"Iris.csv")

        df.head()
```

```
ut[77]:
```

| | Id | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species |
|---|----|---------------|--------------|---------------|--------------|-------------|
| 0 | 1 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 2 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 3 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |

```
In [79]: df.isnull().any()
```

```
Out[79]: Id                False
          SepalLengthCm     False
          SepalWidthCm      False
          PetalLengthCm     False
          PetalWidthCm      False
          Species           False
          dtype: bool
```

```
In [82]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA(n_components=2)
output = lda.fit_transform(X, Y)
outputdf = pd.DataFrame(output, columns=["Linear Discriminant 1", "Linear Discriminant 2"])
finalDf = pd.concat([outputdf, df["Species"]], axis = 1)
```

```
In [83]: finalDf.head()
```

```
Out[83]:
```

| | Linear Discriminant 1 | Linear Discriminant 2 | Species |
|---|-----------------------|-----------------------|-------------|
| 0 | -8.084953 | 0.328454 | Iris-setosa |
| 1 | -7.147163 | -0.755473 | Iris-setosa |
| 2 | -7.511378 | -0.238078 | Iris-setosa |
| 3 | -6.837676 | -0.642885 | Iris-setosa |
| 4 | -8.157814 | 0.540639 | Iris-setosa |

```
In [84]: finalDf.shape
```

```
Out[84]: (150, 3)
```

```
In [128]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

```
In [129]: # Importing the datasets

datasets = pd.read_csv('Salary_Data.csv')
```

```
In [130]: datasets.head()
```

```
Out[130]:
```

| | YearsExperience | Salary |
|---|-----------------|---------|
| 0 | 1.1 | 39343.0 |
| 1 | 1.3 | 46205.0 |
| 2 | 1.5 | 37731.0 |
| 3 | 2.0 | 43525.0 |
| 4 | 2.2 | 39091.0 |

4. Briefly identify the difference between PCA and LDA

Both LDA and PCA rely on linear transformations and aim to maximize the variance in a lower dimension. PCA is an unsupervised learning algorithm while LDA is a supervised learning algorithm. This means that PCA finds directions of maximum variance regardless of class labels while LDA finds directions of maximum class separability.

#PCA

It reduces the features into a smaller subset of orthogonal variables, called principal components – linear combinations of the original variables. The first component

captures the largest variability of the data, while the second captures the second largest, and so on.

#LDA

LDA finds the linear discriminants in order to maximize the variance between the different categories while minimizing the variance within the class.